

# Wir haben die falsche Loss function verwendet?

## Aus der Pytorch Dokumentation:

The input is expected to contain the unnormalized logits for each class

### der unreduzierte Cross Entropy Loss:

$$\ell(x, y) = L = \{l_1, \dots, l_N\}^\top, \quad l_n = -w_{y_n} \log \frac{\exp(x_{n,y_n})}{\sum_{c=1}^C \exp(x_{n,c})} \cdot 1\{y_n \neq \text{ignore\_index}\}$$

Default ist reduction = mean:

$$\ell(x, y) = \begin{cases} \sum_{n=1}^N \frac{1}{\sum_{n=1}^N w_{y_n} \cdot 1\{y_n \neq \text{ignore\_index}\}} l_n, & \text{if reduction} = \text{'mean'}; \\ \sum_{n=1}^N l_n, & \text{if reduction} = \text{'sum'}. \end{cases}$$

Note that this case is equivalent to applying LogSoftmax on an input, followed by NLLLoss

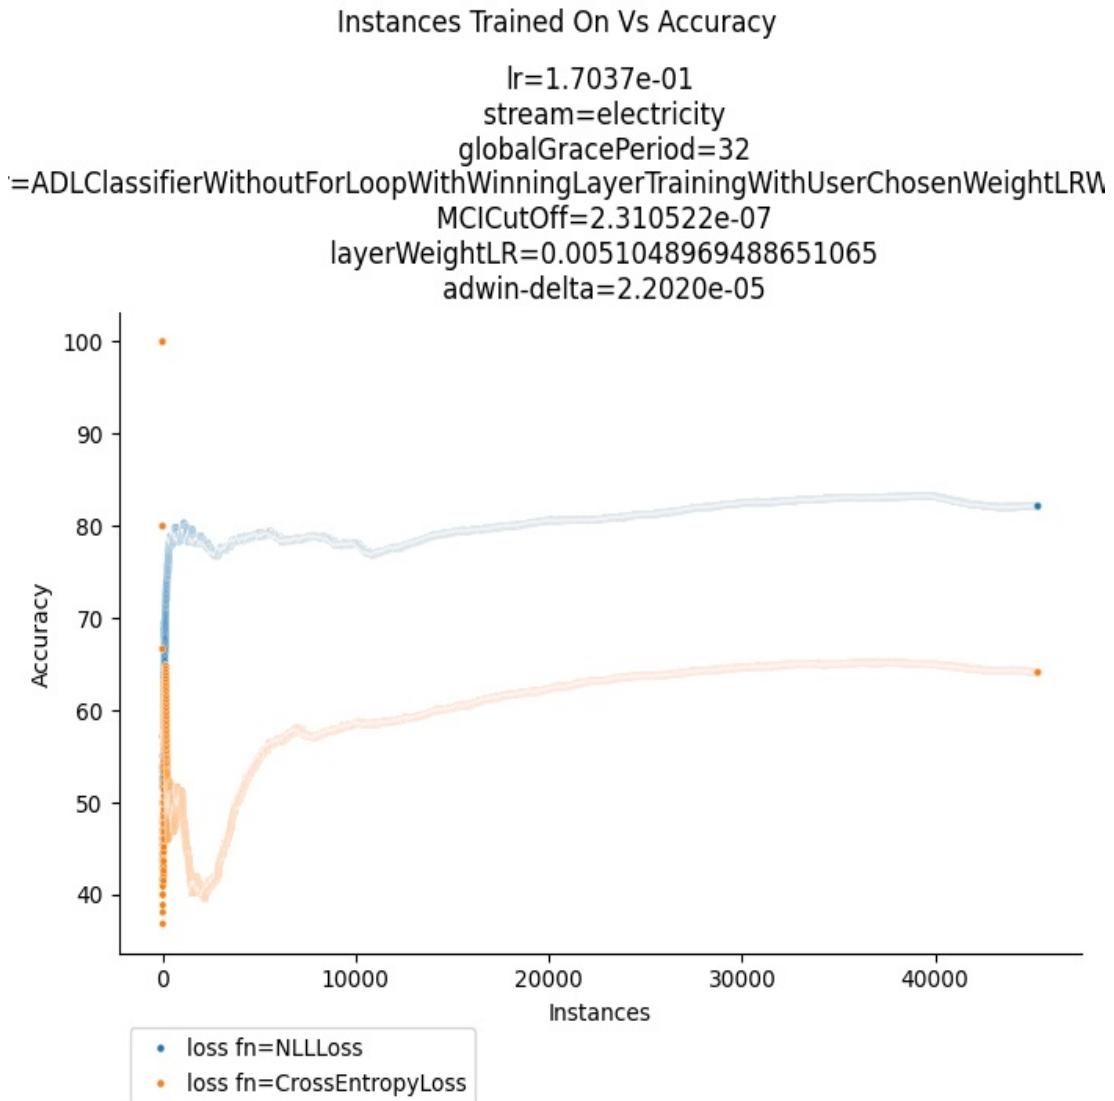
ADL wendet einen Softmax bereits im Forward Pass an, daher ist CrossEntropy vielleicht nicht die geeignete Loss Funktion.

## Vorschlag einer neuen Loss Funktion:

```
import torch
from torch import nn
nr_of_classes, idx_of_true_class = 4, 1

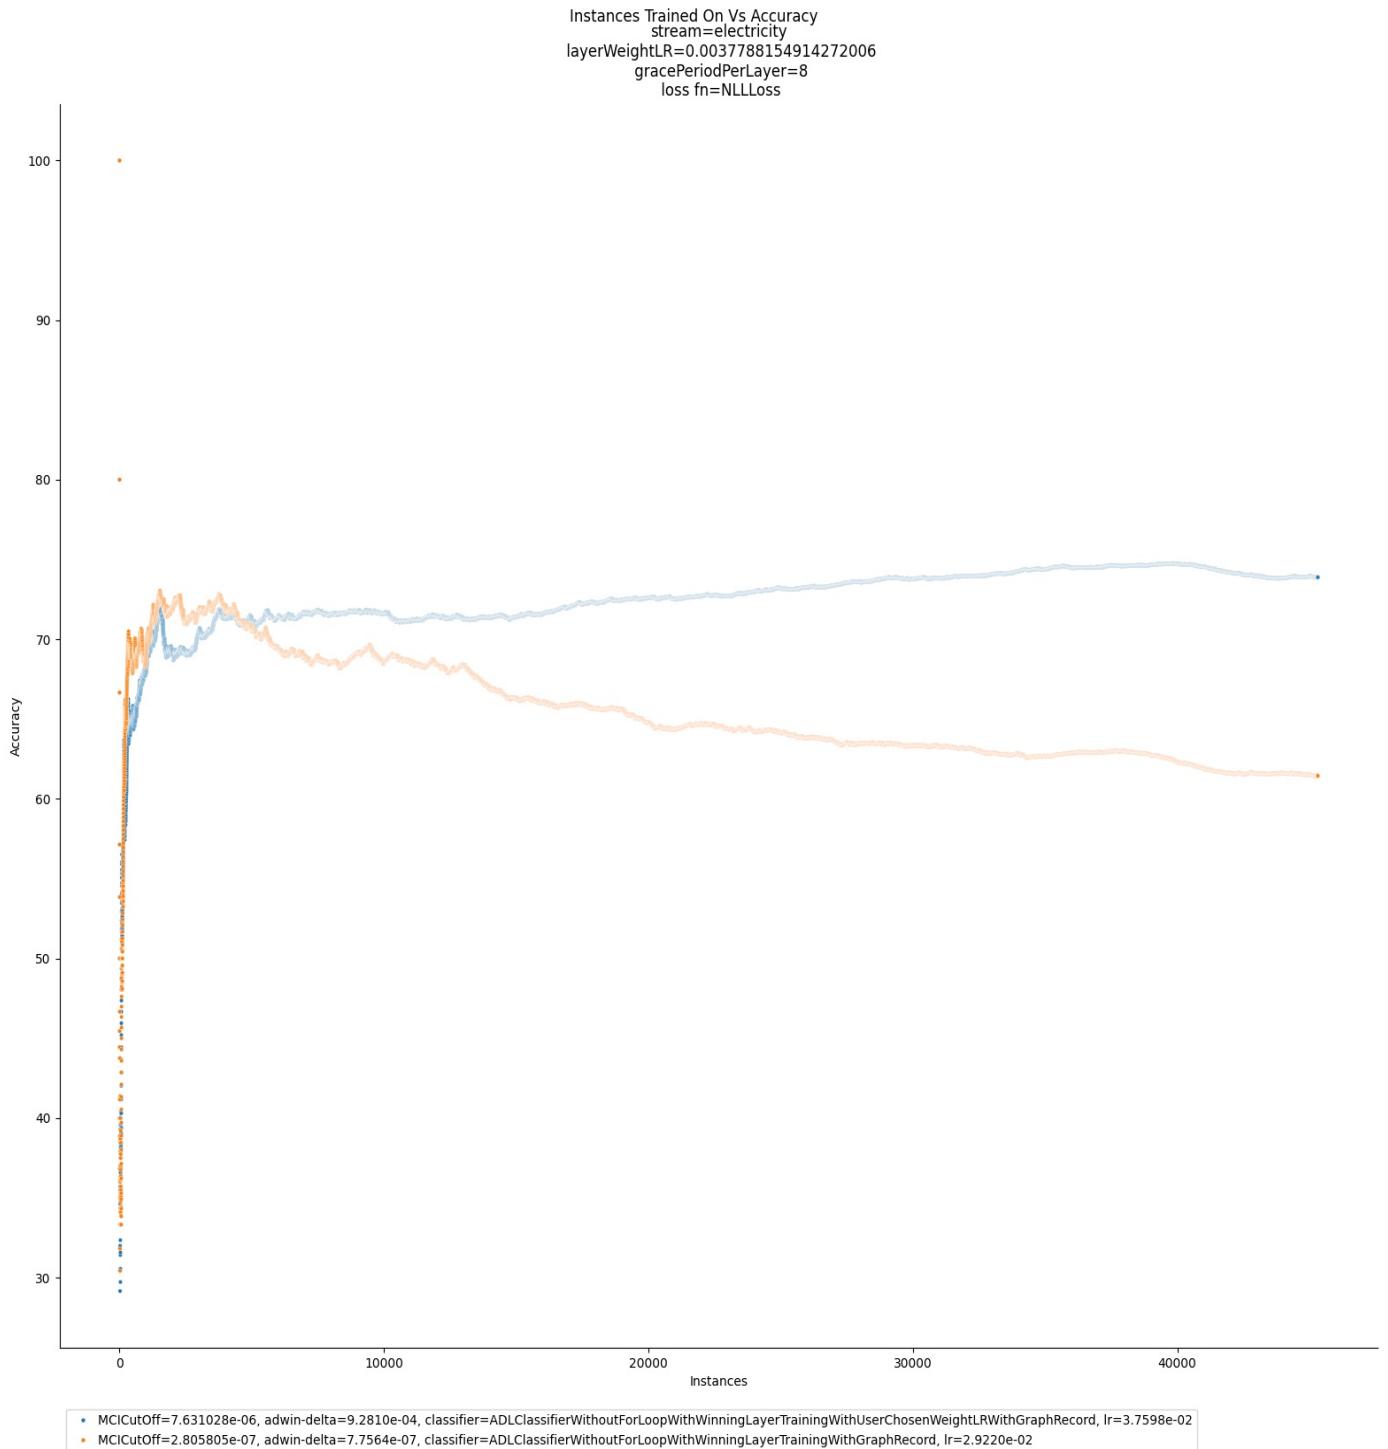
y_pred: torch.Tensor = torch.rand(nr_of_classes)
y_true: torch.Tensor = torch.tensor(idx_of_true_class, dtype=torch.int)
nn.NLLLoss()(torch.log(y_pred), y_true)
```

## Vergleich der alten gegen die neue Loss Funktion für ein Set an Hyperparametern:



**Decoupling von Learning Rate und Weight Correction Factor um LR zu senken?**

**Vergleich best LR Coupled vs Best LR Decoupled**



## Syntetic Streams Build:

### Type of Streams:

#### Type Agrawal SEA details

no drift	y	y	Function 1
one drift	y	y	Function 1 -> abrupt drift @ 5000 -> Function 3
three drifts	y	y	Function 1 -> abrupt drift @ 5000 -> Function 3 -> abrupt drift @ 10000 -> Function 4 -> abrupt drift @ 15000 -> Function 1
drift back and forth	y	y	Function 1 -> abrupt drift @ 5000 -> Function 3 -> abrupt drift @ 10000 -> Function 1

### Results for ADL on Types of Streams

Type	Agrawal	SEA
no drift	53.3%	54.06%
one drift	53.54%	82.00%
three drifts	56.68%	80.95%
drift back and forth	65.71%	81.1%

## Suchraum einschränken:

**1. Versuch: einfach jeweils die drei Parameter nehmen die am besten auf ElectricityTiny performed haben und sie auf Electricity testen:**

Lr:

- LinearLearningRateProgression(initial\_learning\_rate=1, decay\_alpha=0.001)
- ExponentialLearningRateProgression(initial\_learning\_rate=1, decay\_alpha=0.001)
- 5e-1
- 5e-2
- 1e-3

4 Werte weil mir eigentlich alle werte außer den letzten beiden zu hoch sind.

MCI:

- 1e-6
- 1e-7
- 1e-8

in tiny haben wir damit immer so um die 6-10 activen layer am Ende gehabt

adwin-delta:

- 1e-3
- 1e-5
- 1e-7

grace period per layer

- 4
- 8
- 16?
- None

finding aus der isolation: je höher die grace period, desto schlechter das ergebnis und per layer out-performed global nach dem ich global auch anwende #-- es kann sein, dass auf mehr instanzen größere grace periods sinn machen bzw positive sind, weil auf den kleinen datastreams haben wir ja anscheinend nur das problem nicht schnell genug lernen zu können Grace Period None: Sieht sehr lange Trainingszeiten, habe ich dann erstmal zu grace period=1 gemacht.

## Ray Tunes:

kein Grid Search mehr, sondern Probing:

### 1. Suchraum:

- maximal 50000
- frühester abbruch nach = 500
- anzahl stichproben = 500
- learner: (vectorized, winning\_layer, decoupled\_lrs)
- stream: only one stream at a time
- lr: tune.loguniform(1e-4, 5e-1)

- layer\_weight\_learning\_rate: tune.loguniform(1e-4, 5e-1),
- adwin-delta: tune.loguniform(1e-7, 1e-3),
- mci: tune.loguniform(1e-7, 1e-5),
- loss\_fn: NLLLoss
- grace\_period: choice aus: global/layer/none in 4,8,16,32

### Ergebniss des ersten Suchraums:

- lr: 0.17037433308206834,
- layer\_weight\_learning\_rate: 0.0051048969488651065,
- adwin-delta: 2.2019797256079463e-05,
- mci: 2.3105218391180886e-07,
- grace\_period": global, 32

=> 82.15% acc bei 6 hidden, 6 active, und 1502 nodes in hidden layern bei 45000 instancen

### These: min\_runs=500 zu niedrig, bestraft anfänglich langsame lerner

Nachteil von hohem min\_run: suchen dauern sehr lange  
 exemplarisch für Electricity

## 2. Suchraum:

- maximal 50000
- frühester abbruch nach = 500
- anzahl stichproben = 500
- learner: (vectorized, winning\_layer, decoupled\_lrs)
- stream: only one stream at a time
- lr: tune.loguniform(1e-4, **5e-2**), (habe die obere grenze extra niedriger gesetzt um lr zu bekommen die gut sind)
- layer\_weight\_learning\_rate: tune.loguniform(1e-4, **5e-2**),
- adwin-delta: tune.loguniform(1e-7, 1e-3),
- mci: tune.loguniform(1e-7, 1e-5),
- loss\_fn: NLLLoss
- grace\_period: choice aus: global/layer/1 in 4,8,16,32

### Ergebnisse aus 2. Suchraum:

vgl tabelle bei streams beste hyperparameter:

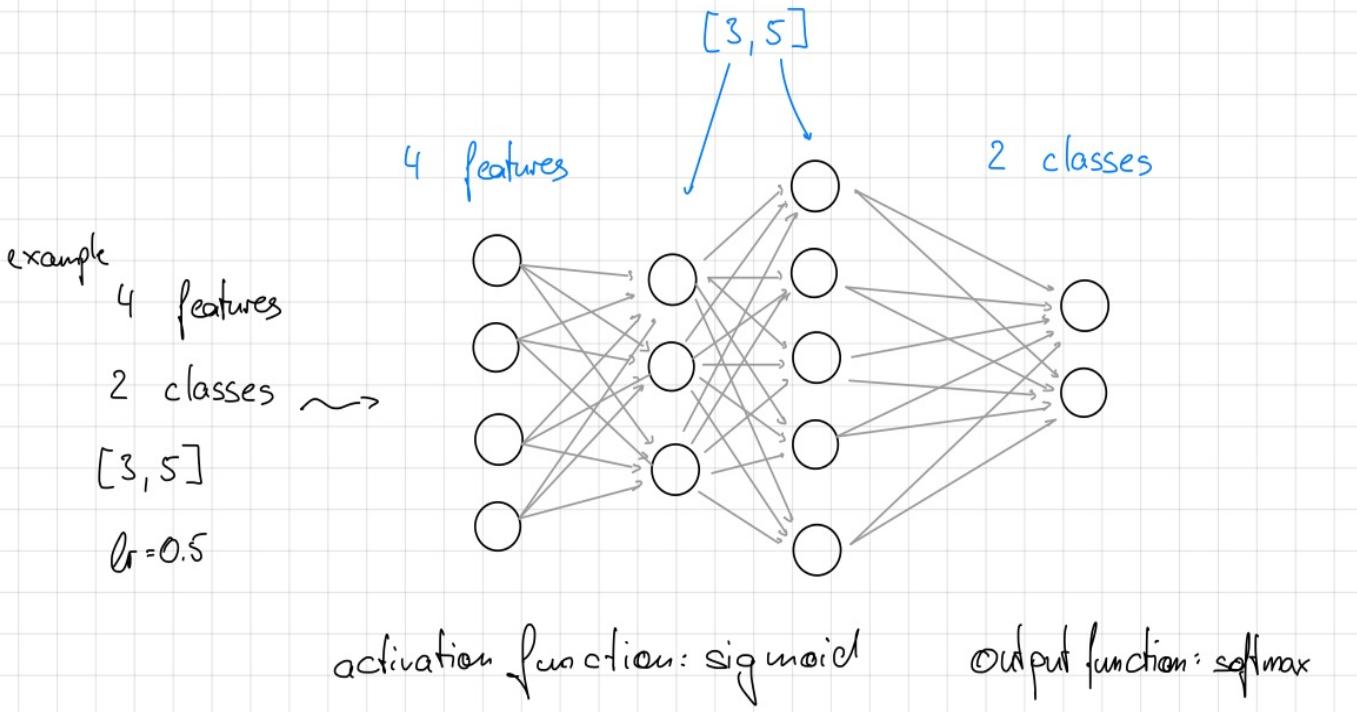
für electricity: python { "learner": [ "vectorized", "winning\_layer", "decoupled\_lrs" ], "stream": "electricity", "lr": 0.03759834306496821, "layer\_weight\_learning\_rate": 0.0037788154914272006, "adwin-delta": 0.0009280964263786628, "mci": 7.631027908752775e-06, "grace\_period": [ 8, "layer\_grace" ], "loss\_fn": "NLLLoss" }

Type	Agrawal					SEA				
	lr	Layer Weight Learning Rate	Adwin Delta	MCI	Grace Period	lr	Layer Weight Learning Rate	Adwin Delta	MCI	Grace Period
no drift	0.001	0.013	3.79e-05	6.8e-07	null	0.001	0.02	4.7e-07	1.2e-06	null
one drift	0.0007	0.0008	1.43e-06	2.48e-07	(8, global)	0.039	0.027	4.14e-07	9.6e-06	(8, global)
three drifts	0.04	0.0001	1.77e-06	8.24e-07	(4, layer)	0.039	0.027	4.14e-07	9.6e-06	(8, global)
drift back and forth	0.023	0.002	0.0002	1.13e-07	(4, global)	0.039	0.027	4.14e-07	9.6e-06	(8, global)

## Comparision Network

## Strukture

input:	nr of features	(size of input vector)
	nr of classes	(size of output vector)
	model structure	how many hidden layers and how big
example	$\ell_r$	



## Results on Electricity

Suche durch den Suchraum: alle Kombinationen an 2er Potenzen an Nodes mit genau so vielen Layern wie das ADL Netzwerk (solange die Anzahl an Layern kleiner als 9 ist, sonst ist Space Komplexität zu groß)  $\leadsto$  python from itertools import product from ray import tune import numpy as np

```
def SimpleDNNSearchSpace(stream_name: str, nr_of_hidden_layers: int = 5, nr_of_neurons: int = 212): """ creates a search space for the SimpleDNN model that has no more than nr_of_hidden_layers many linear layers and in total not more than 2*nr_of_neurons many nodes """
    if nr_of_neurons > 256:
        list_of_possible_neuron_configs = [list(perm) for h in range(1, nr_of_hidden_layers + 1) for perm in product(list(map(int, 2 * np.arange(8, int(np.ceil(np.log2(nr_of_neurons))) + 1))), repeat=h) if np.sum(perm) <= 2 * np.ceil(np.log2(nr_of_neurons)) ]
    else:
        list_of_possible_neuron_configs = [list(perm) for h in range(1, nr_of_hidden_layers + 1) for perm in product(list(map(int, 2 * np.arange(int(np.ceil(np.log2(nr_of_neurons))) + 1))), repeat=h) if np.sum(perm) <= 2 * np.ceil(np.log2(nr_of_neurons)) ]
    return {
        "lr": tune.loguniform(1e-4, 5e-1),
        "model_structure": tune.choice(list_of_possible_neuron_configs),
        stream: tune.grid_search([stream_name])
    }
```
    also nicht jedes mal die gleiche Model Struktur** lr = 0,005  

    Model 1 Layer mit 4096 Nodes Acc: 85.23%

```

## Results for ADL on Types of Streams

Zur Erinnerung:

| Type Agrawal         | SEA    |
|----------------------|--------|
| no drift             | 53.3%  |
| one drift            | 53.54% |
| three drifts         | 56.68% |
| drift back and forth | 65.71% |
|                      | 54.06% |
|                      | 82.00% |
|                      | 80.95% |
|                      | 81.1%  |

## Result for Comparison Network

| Type Agrawal         | SEA     |
|----------------------|---------|
| no drift             | 58.19%  |
| one drift            | 54.53%  |
| three drifts         | 64.31%  |
| drift back and forth | 65.584% |
|                      | 57.57%  |
|                      | 84.25%  |
|                      | 83.94%  |
|                      | 83.61%  |

## Hidden layers Disablen:

### Einfacher Weg:

```
from torch import nn
nr_of_inputs, nr_of_nodes = 3, 4

nn.Linear(nr_of_inputs, nr_of_nodes).requires_grad_(False)
```

Für alle Hidden layer die gelöscht werden. ohne Gradient keine Berechnung von Backward aber im Forward immer noch anwendung der Matrixmultiplikation und der Sigmoidfunction

```
implementiert in : ````python from ADLClassifier.BaseClassifier import ADLClassifier

def disabling_deleted_layers(adl_classifier: type(ADLClassifier)) -> type(ADLClassifier): """
extends an existing ADLClassifier class to disable the gradient calculation for the corresponding
hidden layer if an output layer is deleted :param adl_classifier: the class of ADL Classifier that
should be extended :return: the extended ADLClassifier class """

class DisablingDeletedLayersWrapper(adl_classifier): """ :arg class of ADLClassifier that sets
requires_grad to False for hidden layers whose output layer has been deleted """

    def __str__(self):
        return f"{super().__str__()}WithDisabledDeletedLayers"

    @classmethod
    def name(cls) -> str:
        return f"{adl_classifier.name()}WithDisabledDeletedLayers"

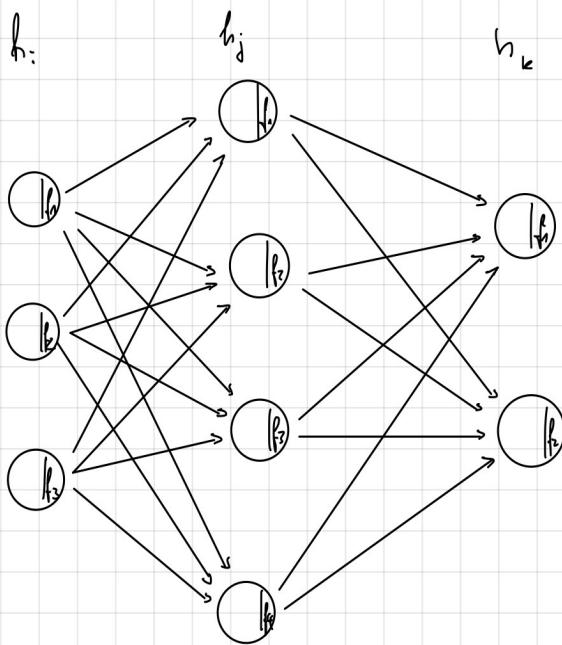
    def _delete_layer(self, layer_index: int) -> bool:
        if super()._delete_layer(layer_index):
            self.model.layers[layer_index].requires_grad_(False)
            return True
        else:
            return False

DisablingDeletedLayersWrapper.name = f"{adl_classifier.name()}WithDisabledDeletedLayers" return
DisablingDeletedLayersWrapper ````
```

### Proposal:

Deletion of hidden Layer Proposal:

Skizze: Vorher



$$f \left( h_k \left( f \left( h_j \left( f \left( h_i(x) \right) \right) \right) \right) \right)$$

in example :  $n = 3$

$m = 4$

$o = 2$

$n \times |x|$

$$h_i \in \mathbb{R}^n$$

$$h_j \in \mathbb{R}^{m \times n}$$

$$h_k \in \mathbb{R}^{o \times m}$$

$$h_l' = f \circ h_j \circ f \circ h_k$$

Problem:  $f$  is non-linear      Proposal :

$\Rightarrow f \circ h_j \circ f \circ h_k$  cannot be reduced

to a single layer

$$h_l'' = \underbrace{f \circ h_j \circ h_k}_{\text{one matrix after multiplication}}$$

→ one layer  
+ 1 sigmoid

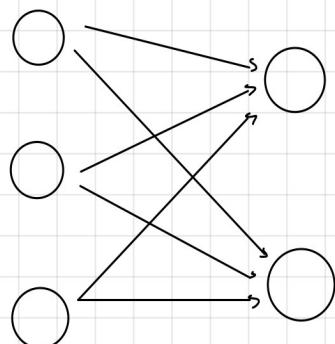
Nachher:  $h_i$

$$h_k'$$

$$f \left( h_k' \left( f \left( h_i(x) \right) \right) \right)$$

With Bias:

$$\begin{aligned} H_2 \cdot (H_1 \cdot x + b_1) + b_2 \\ = H_2 H_1 x + H_2 b_1 + b_2 \end{aligned}$$



By merging with the layer after the deleted  
we don't change the output size of newly merged  
compared to the remaining one  
→ we need to change the output layer

implementiert in: ```python from ADLClassifier import ADLClassifier

```
def delete_deleted_layers(adl_classifier: type(ADLClassifier)) -> type(ADLClassifier):
    """ extends an existing ADLClassifier class to disable the gradient calculation for the corresponding hidden layer if an output layer is deleted
    :param adl_classifier: the class of ADL Classifier that should
```

```

be extended :return: the extended ADLClassifier class """
class DeleteDeletedLayersWrapper(adl_classifier): """ :arg class of ADLClassifier that sets
requires_grad to False for hidden layers whose output layer has been deleted """
def __str__(self):
    return f"{super().__str__()}WithDeleteDeletedLayers"

@classmethod
def name(cls) -> str:
    return f"{adl_classifier.name()}WithDeleteDeletedLayers"

def _delete_layer(self, layer_index: int) -> bool:
    # not exactly the same output, as we remove a sigmoid function between both layers in the forward s
    if super()._delete_layer(layer_index):
        self.model.delete_hidden_layer(layer_index)
        return True
    else:
        return False

DeleteDeletedLayersWrapper.name = f"{adl_classifier.name()}WithDeleteDeletedLayers" return
DeleteDeletedLayersWrapper sowie: python from typing import List, Dict, Optional, Tuple
import numpy as np import torch from torch import nn
class AutoDeepLearner(nn.Module):

def delete_hidden_layer(self, layer_index: int) -> None:
    """
    deletes a hidden layer whose output layer was deleted beforehand
    :param layer_index: the index of the hidden layer to delete
    """
    # make sure that there is another layer
    assert len(self.layers) > 1, "there needs to be at least another layer to delete a hidden layer"
    assert not self.output_layer_with_index_exists(layer_index), "cannot delete a hidden layer of an act

    layer_to_delete: nn.Module = self.layers[layer_index]
    # merge two layers and take the place of the second layer:
    if layer_index == len(self.layers) - 1:
        # if last layer is deleted just delete the layer, no need to merge anything
        pass
    else:
        # if not first layer is deleted merge with layer in after it
        layer_to_merge_with: nn.Module = self.layers[layer_index + 1]
        new_weight = nn.Parameter(layer_to_merge_with.weight.matmul(layer_to_delete.weight))
        new_bias = nn.Parameter(layer_to_merge_with.bias + layer_to_merge_with.weight.matmul(layer_to_de
        new_layer = nn.Linear(layer_to_delete.in_features, layer_to_merge_with.out_features)
        new_layer.weight = new_weight
        new_layer.bias = new_bias
        self.layers[layer_index + 1] = new_layer

    # remove now merged layer
    self.layers.pop(layer_index)

    # update all keys of all weights, all output keys and all weight_correction_factors
    # for all active layers that follow the deleted layer:
    for active_key in self.active_layer_keys()[self.active_layer_keys() > layer_index].detach().numpy():
        self.__set_output_layer(active_key - 1, self.__pop_output_layer(active_key))

```

```
self.__set_voting_weight(active_key - 1, self.__pop_voting_weight(active_key))
self.__set_weight_correction_factor(active_key - 1, self.__pop_weight_correction_factor(active_k
```

## **Ergebnisse des disablen von hidden layern:**

keine zeit mehr für runs gehabt

### **Accuracy Changes**

### **Emission Changes**

### **Notizen:**

1. Wenn Concept Change passiert macht eine Learning Rate Progression nur Sinn wenn sie dann wieder von vorne beginnt  
-> Future Work (nach dem 20.03.)
2. Future Work: Write Capymoa classifier that runs the Matlab Implementation (for benchmarking reasons)